

Large language models instantiate evolutionarily robust strategies of cooperation

Saptarshi Pal^{1*}, Abhishek Mallela², Lenz Pracher³, Chiyu Wei², Feng Fu^{2,4}, Santiago Schnell^{2,4,5}, and Martin A. Nowak^{1,6}

¹Department of Mathematics, Harvard University, Cambridge, MA 02138, USA

²Department of Mathematics, Dartmouth College, Hanover, NH 03755, USA

³Arnold Sommerfeld Center for Theoretical Physics, Ludwig-Maximilians-Universität, München 80539, Germany

⁴Department of Biomedical Data Sciences, Geisel School of Medicine, Hanover, NH 03755, USA

⁵Department of Biochemistry and Cell Biology, Dartmouth College, Hanover, NH 03755, USA

⁶Department of Organismic and Evolutionary Biology, Harvard University, Cambridge, MA 02138, USA

*To whom correspondence should be addressed: Email: spal@math.harvard.edu

Edited By Attila Szolnoki

Abstract

Large language models (LLMs) are increasingly deployed to support human decision-making. This use of LLMs has concerning implications, especially when their prescriptions affect the welfare of others. To gauge how LLMs make social decisions, we explore whether five leading models produce sensible strategies in the repeated prisoner's dilemma, which is the main metaphor of reciprocal cooperation. First, we measure the propensity of LLMs to cooperate in a neutral setting, without using language reminiscent of how this game is usually presented. We record to what extent LLMs implement Nash equilibria or other well-known strategy classes. Thereafter, we explore how LLMs adapt their strategies to changes in parameter values. We vary the game's stopping probability, the payoff values, and if the total number of rounds is commonly known. We also study the effect of different framings. In each case, we test whether the adaptations of the LLMs are in line with basic intuition, theoretical predictions of evolutionary game theory, and experimental evidence of human participants. While all LLMs perform well in many of the tasks, none of them exhibit full consistency over all tasks. We also conduct tournaments between the inferred LLM strategies and study direct interaction between LLMs in games over ten rounds with known or unknown last round. Our experiments shed light on how current LLMs instantiate reciprocal cooperation.

Significance statement

Large language models (LLMs), such as Claude and GPT, are increasingly becoming a part of everyday decision-making. As they become more pervasive, it becomes important to understand whether they support mutually beneficial interactions. We investigate this question by studying how five leading LLMs behave in the repeated Prisoner's Dilemma, a canonical model of reciprocal cooperation. We estimate the "effective strategies" that LLMs implement when given limited information about past outcomes. We find that their behavior is generally cooperative and generous, and often evolutionarily robust. However, LLMs occasionally respond in ways that diverge from theoretical predictions and intuitive expectations. This highlights limitations in their ability to consistently sustain cooperative behavior.

Introduction

AI and large language models (LLMs) are entering into all areas of human activity (1–8). As they become more powerful and more widely used, it is important to gauge their propensity to cooperate or defect in strategic situations (4, 9–17). Studying

cooperative decision-making is particularly relevant in social dilemmas, in which there is a conflict between cooperation and defection. Cooperation maximizes the payoff of a community, while defection tempts individuals to exploit short-term benefits (18–23). A classical setting is the prisoner's dilemma (24–26), where two players have to choose simultaneously between

Received: March 13, 2026. Accepted: June 1, 2026

© The Author(s) 2026. Published by Oxford University Press on behalf of National Academy of Sciences.

This is an Open Access article distributed under the terms of the Creative Commons Attribution-NonCommercial License (<https://creativecommons.org/licenses/by-nc/4.0/>), which permits non-commercial re-use, distribution, and reproduction in any medium, provided the original work is properly cited. For commercial re-use, please contact reprints@oup.com for reprints and translation rights for reprints. All other permissions can be obtained through our RightsLink service via the Permissions link on the article page on our site—for further information please contact journals.permissions@oup.com.

Box 1 Prompt used for the first experiment You are in an interaction with another agent.

Your interaction is over several rounds.

In each round, you have a choice between two options, L or R.

The other agent also has a choice between L or R.

The two of you will decide without knowing what the other will choose.

Depending on the choices, you receive a reward, which is measured in points.

If both of you choose L, both of you get $[a_{LL}]$ points.

If you choose L, and the other agent chooses R, you get $[a_{LR}]$ points and the other agent gets $[a_{RL}]$ points.

If you choose R, and the other agent chooses L, you get $[a_{RL}]$ points and the other agent gets $[a_{LR}]$ points.

If both of you choose R, both of you get $[a_{RR}]$ point(s).

Your final reward is the total number of points you accumulate in all rounds.

Here is your goal: No instruction.

In the previous round, you chose $[x]$, the other agent chose $[y]$.

Therefore, you got $[a_{xy}]$ point(s) and the other agent got $[a_{yx}]$ point(s).

This is a new round.

Do you choose L or R? Give only the character as output. Give no explanation.

cooperation, C, and defection, D. If both cooperate, they get a higher payoff than if both defect. But if one of them defects while the other cooperates, the defector gets the highest payoff of the game and the cooperator, the lowest. This game is described by the payoff matrix

	C	D
C	a_{CC}	a_{CD}
D	a_{DC}	a_{DD}

(1)

The entries of the matrix indicate the payoffs for the row player. The game is a prisoner's dilemma if $a_{DC} > a_{CC} > a_{DD} > a_{CD}$. This ranking implies that defection dominates cooperation in the one-shot game: no matter what the co-player will do, the focal player maximizes her payoff by defecting. Defection is the only Nash equilibrium (27): if both players defect neither one can unilaterally change to cooperation and thereby increase his payoff.

But even selfish players may find it worthwhile to cooperate if the game is played for several rounds (28). In the repeated prisoner's dilemma, cooperation can prevail by a mechanism called direct reciprocity (20): If I cooperate now, you may cooperate later. If I defect now, you may defect later. Thus, the shadow of the future provides an incentive to cooperate in the present. A typical setting for the repeated game is obtained by using a probability, w , that the interaction will end after each round. A choice of $w = 1$ recovers the one-shot game explained above. The converse limit, $w \rightarrow 0$, describes the infinitely repeated game without discounting. There is a large literature on evolution of cooperation by direct reciprocity (29–41). Prominent strategies for playing the repeated prisoner's dilemma include tit-for-tat (29), generous tit-for-tat (31), win-stay, lose-shift (32), GRIM (42), and Forgiver (43, 44). See Methods for details.

We examine the behavior of five leading LLMs in the repeated prisoner's dilemma. They are: (i) claude-sonnet-4 by Anthropic, (ii) gemini-2.5-pro by Google DeepMind, (iii) gpt-4o and (iv) gpt-5 by OpenAI, and (v) llama-3.3-70b by Meta. We abbreviate these models as Claude, Gemini, Gpt-4o, Gpt-5 and Llama

(Fig. 1A). In all our experiments, the parameter settings of the LLMs are chosen to place them into regimes that are known to produce relatively deterministic outputs (see Methods for the parameterization).


Prior work has begun to explore how LLMs play repeated games (45–54). In those studies, LLMs tend to be pitted against other LLMs or against simple computer programs. Those efforts provide valuable insights into the dynamics of LLM play (which we review in detail in the [Supplementary material](#)). However, they primarily describe realized outcomes of individual experiments—rather than the underlying LLM strategies that generate these outcomes. This limitation makes it difficult to connect the observed LLM behavior to formal solution concepts in evolutionary game theory. In contrast, we infer LLM strategies directly, by asking LLMs to react to every possible outcome of the last one (or two) rounds. While this change in approach may seem innocuous, it has dramatic consequences. It allows us to ask how LLMs would fare against any other (hitherto unobserved) opponent strategies. Moreover, by eliciting the LLMs' strategies directly, we can compare their behavior to the considerable literature on evolution of cooperation (29–41). We can directly analyze the properties of those strategies using concepts from (evolutionary) game theory (55–57).

We use this approach to ask whether current LLMs produce sensible strategies of direct reciprocity. First, we describe which behaviors LLMs instantiate in a neutral setting. We avoid any language traditionally used to present the repeated prisoner's dilemma. In particular, we examine whether the generated strategies are Nash equilibrium strategies (27), and whether they fall into the formal classes of “partners” or “rivals” (58), which characterize the degree to which a strategy is generous or competitive. There is a correspondence between Nash equilibrium and evolutionarily robust strategies. A strategy is evolutionarily robust if no mutant strategy starting as a single invader is favored by natural selection to replace it. In the limit of infinite population size and strong selection, the concepts of Nash equilibrium and evolutionary robustness coincide (38). Partner strategies are therefore evolutionarily robust cooperative strategies in that limit.


A Description of the considered LLMs, and the five events relevant for memory-1 strategies.

LLM model name	Shorthand	Scenario	Description
1. claude-sonnet-4	CLAUDE	R1	Round one.
2. gemini-2.5-pro	GEMINI	LL	In the previous round both played L.
3. gpt-4o	GPT-4O	LR	In the previous round, agent played L, co-player played R.
4. gpt-5	GPT-5	RL	In the previous round, agent played R, co-player played L.
5. llama-3.3-70b	LLAMA	RR	In the previous round, both played R.

B Cooperation frequency by LLMs in the five memory-1 scenarios in fifty independent trials.



C Range of expected payoffs in games between inferred LLM strategies and arbitrary opponents.



□ : Constitutes Nash equilibrium during self-play
(Percentage indicates the share of random memory-1 opponents that beat strategy's self-payoff).

P/R : Classifies as Partner/Rival.

Figure 1 How do LLMs respond to memory-1 scenarios of the repeated prisoner's dilemma? A) We consider five state-of-the-art LLMs—claude-sonnet-4 by Anthropic, gemini-2.5-pro by Google DeepMind, Gpt-4o and Gpt-5 by OpenAI and llama-3.3-70b by Meta AI. We refer to these models Claude, Gemini, Gpt-4o, Gpt-5, and Llama. We present five repeated prisoner's dilemma scenarios to the LLMs and request their choices. The five scenarios elicit the LLM's response in the first round, and after the four possible outcomes of the previous round (see Box 1 for the prompts). B) We report how often the five LLMs choose the cooperative action "L" in each scenario across 50 independent trials. See Tables S3–S7 for the data. C) We show the range of expected per-round payoffs possible when inferred LLM memory-1 strategies play arbitrary opponents. We show these ranges for games with infinitely many rounds (top row) and games with 100 rounds in expectation (bottom row). We indicate whether an LLM strategy gives rise to a Nash equilibrium with a yellow square; otherwise, we report the percentage of randomly sampled memory-1 opponents (out of 10^6) that achieve a higher payoff against it than it earns against itself. With "P" or "R" we denote if the LLM strategy is a partner or a rival.

In a second step, we record how LLMs adapt their strategies as we vary several key parameters, including the game's stopping probability, payoff values, the number of past rounds the LLM is instructed to keep in memory, and whether or not players are aware that the current round is the last. All parameters are known to affect both theoretical predictions (36, 57, 59, 60) and human behavior (61, 62). We explore whether changes in these parameters elicit similar adaptations among LLMs. Interestingly, all of the LLMs show sensible adaptations in some of the tasks, but none of them exhibit full consistency across all of the tasks.

In a last step, we explore the robustness of the LLM output by using different framings of the decision task (either eliciting more competitive or more cooperative behavior). Moreover, we conduct several round-robin tournaments (29, 63). These analyses illustrate to which extent LLM behavior can be steered into certain directions, and how LLMs fare when competing against each other. Overall, this research provides an important glimpse into how current generations of LLMs instantiate reciprocal cooperation. In view of the recent progress in agentic AI, such artificially generated social behaviors are increasingly likely to produce spillovers in the world we live in.

The first experiment

In our first experiment, we test how the five LLMs choose between cooperation and defection in a repeated interaction. We use the prompt shown in Box 1 to obtain their choice, depending on the outcome of the previous round. We use a slightly different prompt to obtain their choice for the first round (see Methods). The wording is designed to avoid ambiguity. The labels of the two options, L and R, are used as neutral symbols for cooperation and defection. For the payoff values in this first experiment, we use the standard parameters (29), $(a_{LL}, a_{LR}, a_{RL}, a_{RR}) = (3, 0, 5, 1)$. For those values the game is a prisoner's dilemma, and mutual cooperation, LL, receives a higher payoff than alternating between LR and RL. We ask all five LLMs 50 times for each of five decisions, namely: what to do in the first round, and what to do if the previous round was LL, LR, RL, or RR. Thus, we record the propensity of the LLM to cooperate (i) in the first round, (ii) after mutual cooperation (LL), (iii) after being exploited (LR), (iv) after exploiting the co-player (RL), and (v) after mutual defection. If the LLMs choose to act consistently with Nash equilibrium play, we would expect them to generate strategies that either fully cooperate, fully defect, alternate between cooperation or defection when playing against a copy of themselves, or that they implement so-called equalizers (59, 60).

The results of our experiments are shown in Fig. 1B. Four of the five LLMs cooperate fully in the first round. Only Gemini is suspicious: it cooperates 41 times in 50 trials. Claude implements the strategy Forgiver (43, 44): it cooperates fully after mutual cooperation, after having exploited the co-player and after mutual defection; it always defects if it was exploited. Gemini uses a stochastic variation of WSLS (32): it cooperates fully after mutual cooperation, it never cooperates after having been exploited; it cooperates 1/50 times after exploiting the co-player; it cooperates 21/50 times after mutual defection. Gpt-4o uses GRIM (56): it fully cooperates after mutual cooperation and fully defects otherwise. Gpt-5 uses a stochastic strategy that is between WSLS and Forgiver; after exploiting the co-player it cooperates

27/50 times. Llama uses a version of generous-tit-for-tat (31, 64): it always cooperates after mutual cooperation; it cooperates 13/50 times after being exploited; it cooperates 49/50 times after exploiting the co-player; it cooperates 28/50 times after mutual defection. See [Supplementary material](#) for all data.

We can use these data to infer memory-1 strategies (see Tables S11–S20). A memory-1 strategy is given by five parameters, $(p_0, p_{LL}, p_{LR}, p_{RL}, p_{RR})$. These parameters represent the probabilities to cooperate in the opening round, p_0 , and the probabilities to cooperate after each of the four outcomes of the previous round. Based on the inferred memory-1 strategy, we can also compute whether an LLM's strategy is a Nash equilibrium (27), a partner, or a rival (58). A strategy is a Nash equilibrium if the expected payoff it earns against itself is at least as high as the expected payoff any opponent earns from it. A strategy is a *partner* if (i) it is a Nash equilibrium and (ii) if it fully cooperates with a copy of itself. A strategy is a *rival* if its own payoff is greater than or equal to the opponent's payoff for any opponent. Partner strategies are a subset of Nash equilibria, but rival strategies are not. Whether a given strategy is a member of any of these strategy classes depends on the payoff values, $(a_{LL}, a_{LR}, a_{RL}, a_{RR})$, and the termination probability, w , see Ref. (57). Most memory-1 strategies are in neither of the three sets (57).

Based on these definitions, we explore whether any of the LLM strategies is a Nash equilibrium, a partner, or a rival. For each strategy, we also calculate the range of expected payoffs that are attained in an encounter between this strategy and an arbitrary opponent. The result is shown in Fig. 1C, using either a stopping probability of $w = 0$ or $w = 10^{-2}$. In the infinitely repeated game ($w = 0$), all five LLMs produce partner strategies. In particular, all strategies are Nash equilibria. However, once we consider long but not infinitely long games ($w = 10^{-2}$), four of the five LLM strategies remain fully cooperative, but only two of them remain Nash equilibria (the strategies generated by Gpt-4o and Llama).

Dependence on parameter changes

Impact of the stopping probability

The previous experiments establish a baseline scenario. They show that in a commonly used parameter domain, LLMs aim to establish cooperation. In a next step, we explore how LLM behaviors change as we systematically vary the parameters of the game. We begin by varying the stopping probability w . Theory and experimental work both predict that larger values of w make it more difficult to sustain cooperation. After all, a larger stopping probability reduces “the shadow of the future” (57, 62). We would thus expect that larger values of w render the LLM strategies less tolerant (57). That is, the cooperation probabilities after unilateral or mutual defection ought to become smaller in a monotonic fashion. Moreover, above a certain value of w , an LLM may stop to cooperate altogether.

In our original experiment, we only told the LLMs that the interaction is over “several rounds”. Now we add the following sentence: “After each round the interaction ends with probability $[w]$ ”. For exact prompt details, see **Methods**. We use the following four values $w = 0.01, 0.1, 0.2, 0.5$. The results are shown in Fig. 2A. Curiously, Claude and Llama are unaffected by the value of w . Claude always opts for Forgiver (even though the strategy

ceases to be a Nash equilibrium for $w > 0$). Similarly, Llama always chooses GRIM. Gpt-4o produces behavior that is in between GRIM and WSLS. Interestingly, it exhibits some nonmonotonic variation in p_{RR} , its cooperation probability after mutual defection. Gemini reduces its probability to cooperate in the first move and after mutual cooperation as w increases. Gpt-5 reduces its propensity to cooperate after exploiting the coplayer and after mutual defection. These latter two LLMs adapt their strategies in ways that are most consistent with our expectation of monotonically decreasing cooperation propensities. Interestingly, however, for $w > 0$, neither Gemini nor Gpt-5 produce strategies that are consistent with Nash equilibrium play (see Table S27). Instead, Llama produces a partner strategy for all values of w , while Gpt-4o's strategy is a partner for all w except $w = 0.5$. (Recall that every partner strategy is a Nash equilibrium by definition.)

Changing the payoff matrix

As another key parameter, we change the players' incentives to cooperate. We conducted 11 additional experiments using the following payoff values: 10 for mutual cooperation; 0 for cooperation versus defection; $10 + x$ for defection versus cooperation; and x for mutual defection. We vary the parameter $x = 0, 1, 2, \dots, 10$. Increasing x means the cost for cooperation increases and the temptation to defect becomes larger. For $x = 0$, cooperation has no cost at all. For $x = 10$, the cost of cooperation is as large as the benefit. For all values of x in between 0 and 10, the game is a prisoner's dilemma. The larger this cost x , the more difficult cooperation becomes to sustain. Thus, we would expect that with increasing costs, the cooperation propensities of LLMs decrease again in a monotonic fashion, see Refs. (38, 57).

The results of the LLM experiments are shown in Fig. 2B. Claude uses its favorite strategy Forgiver for all values of $x < 10$; for $x = 10$ it suddenly uses TFT. Gemini chooses WSLS for small costs, but switches to suspicious GRIM for large costs. Gpt-4o uses WSLS for small costs and switches to GRIM for large costs. Gpt-5 uses Forgiver for small costs and reduces its cooperation propensity for larger costs; for very large costs it approaches "always defect" (ALLD). Llama uses "always cooperate" (ALLC) for zero cost. For larger costs, it maintains full cooperation in the first round and after mutual cooperation, but it reduces its cooperation propensity after mutual defection. Overall, these changes are largely in line with our expectations. However, again we observe some remarkable nonmonotonicities in some of the cooperation propensities (especially for Llama, Fig. 2B).

Again, we analyze whether the inferred strategies are Nash equilibria, partners, or rivals for $w = 0$ and $w = 10^{-2}$. The results are summarized in Table S29. Firstly, we observe that the chosen strategies are predominantly partners and not rivals. Secondly, strategies are more likely to be partners when x is sufficiently low. In particular, save for a few exceptions, these strategies are only partners for $x \leq 5$. Only Gpt-4o's strategy classifies as a partner for most x .

Overall, the above results suggest that all LLMs perform well in some of the tasks. However, none of the LLMs exhibit sensible behavior throughout. For example, Claude seems to ignore changes in the game's stopping probability or in the incentives to cooperate (Fig. 2), whereas the cooperation probabilities of

Gpt-4o and Llama change nonmonotonically (based on 50 independent API calls).

Memory-2 decisions of LLMs

While many publications explore optimal behavior among memory-1 strategies, the literature on higher memory is comparably scarce (36, 37, 66–68). This raises the question whether LLMs—which are trained on the published literature—are able to generate reasonable strategies in these lesser explored domains. To explore that question we have conducted experiments where we study decisions of LLMs given the outcome of the previous two rounds. For the exact prompts, see **Methods**. The results are shown in Fig. 3. For the first two rounds, the results coincide with the original treatment (Fig. 1). For the subsequent rounds, we find the following. If the previous round was mutual cooperation (LL) then Gemini, Gpt-4o, Gpt-5, and Llama cooperate fully regardless of the outcome two rounds ago; only Claude looks further back and chooses to defect if it was exploited two rounds ago. If the agent was exploited in the previous round (LR), then Claude, Gemini, and Gpt-4o fully defect regardless of what happened two rounds ago, Gpt-5 is willing to cooperate if it exploited the co-player two rounds ago; but Llama is willing to cooperate unless it was exploited twice in a row. In general, Llama is very cooperative in this setting. Analyzing the resulting memory-2 strategies, we find that Claude, Gpt-4o, and Gpt-5 are partners, whereas Gemini and Llama are neither Nash equilibria nor rivals.

Repeated games over 10 rounds

The previous experiments assume that players are not aware of how many rounds they will play. This is a standard assumption made to preclude end game effects (55): Once the duration of the game is known, rational players should defect in the last round; but once players expect their opponents to defect in the last round, they should already defect the round before, and so on eg (61, 71). Depending on how far-sighted LLMs are—or to which extent they expect their opponent to be far-sighted—cooperation may thus unravel completely. To explore whether LLMs engage in this kind of backward induction, we study a repeated prisoner's dilemma with ten rounds. We performed two treatments. In one treatment, we tell the LLMs that the game lasts at least 10 rounds. In the other treatment, we specify that we play *exactly* 10 rounds.

Each LLM plays against all five LLMs. This results in a total of 15 distinct encounters. As the games proceed, we remind the LLMs of the cumulative history of the game until the present round. For the exact prompts, see **Methods**. In the first treatment, we observe full cooperation in all games, which is consistent with the data from the original experiment. In the second treatment, all LLMs except Gpt-5 cooperate in every round of every game, suggesting that none of them exercise backward induction. Only Gpt-5 consistently defects in the final round across all its games, exhibiting a clear end-game effect, while cooperating in all preceding rounds. This does not constitute full backward induction, which would predict defection in every round, not just the last.

Note that all LLMs except Gpt-5 fail to defect in the last round when the game has a known ending. This failure to exploit the opponent is remarkable, because defection in the last round is

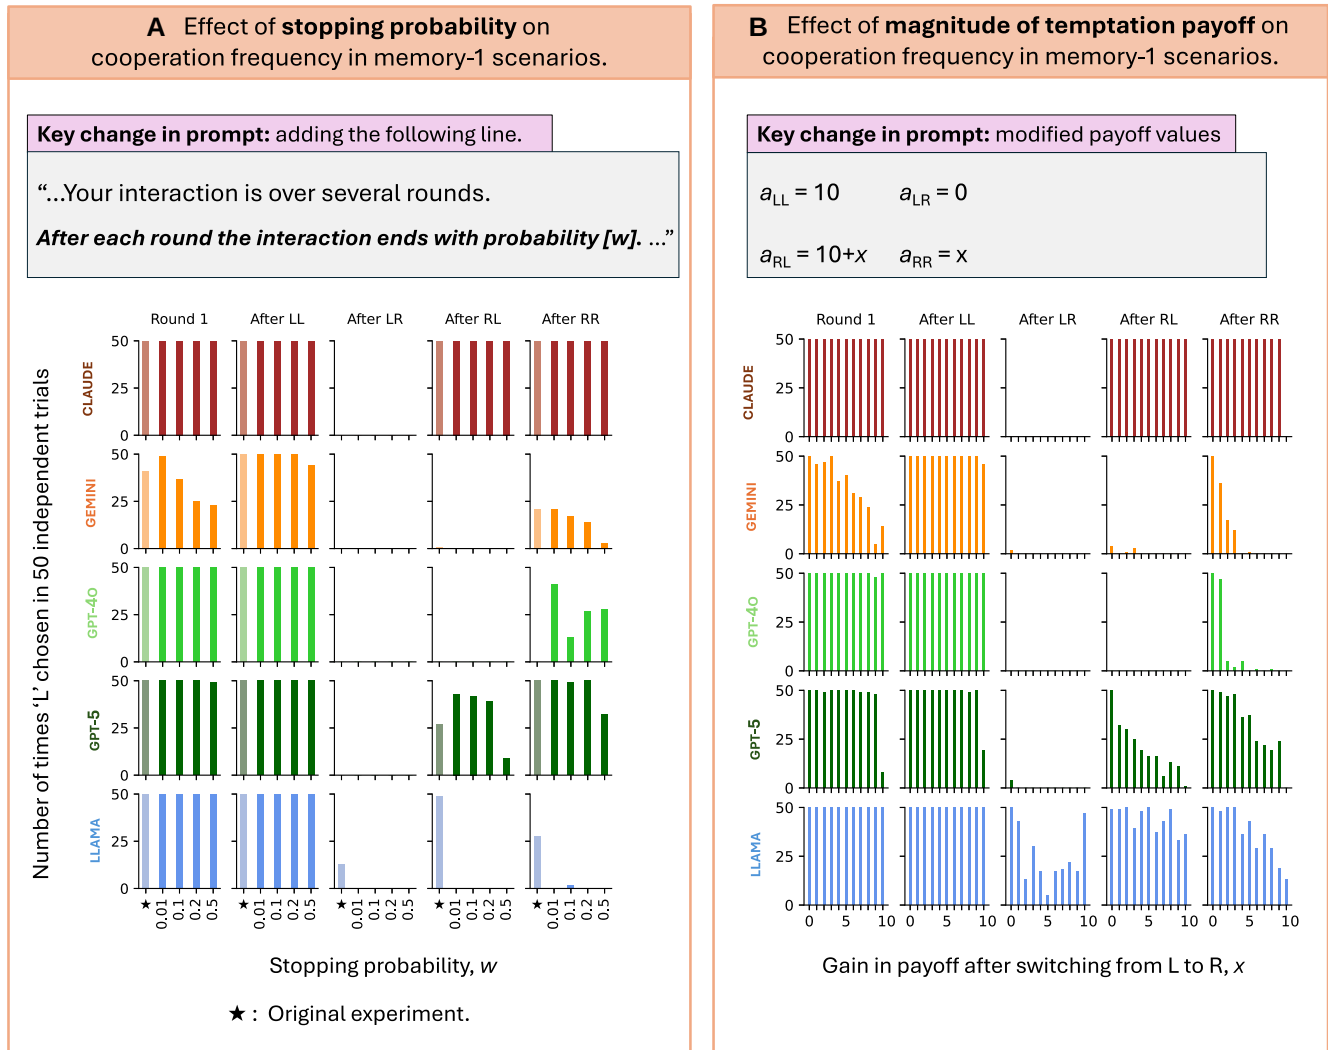


Figure 2 How does conditional cooperation by LLMs vary with the game parameters? We vary two key game parameters: the interaction’s stopping probability (A) and the game payoffs (B). For the first, we add a single line to the original prompt in Box 1: “After each round the interaction ends with probability $[w]$.” We vary w and record how often the five LLMs choose the cooperative action “L” in the five memory-1 scenarios across 50 independent trials. The payoff parameters match those used in the previous figure, $(a_{LL}, a_{LR}, a_{RL}, a_{RR}) = (3, 0, 5, 1)$. For the second, we use the special payoff values $(a_{LL}, a_{LR}, a_{RL}, a_{RR}) = (10, 0, 10 + x, x)$, representing games with the “equal gains from switching” property (65). Here, x is the gain in payoff when unilaterally shifting from the cooperative action “L” to the defective action “R.” At $x = 0$, a player’s payoff is independent of its own action. At $x = 10$, cooperation ceases to be socially optimal since mutual defection is as beneficial as mutual cooperation. We vary x and plot how often the five LLMs cooperate in the five memory-1 scenarios. For more data and information on whether the inferred strategies form Nash equilibria or act as partners or rivals, see [Section 4.2 of the Supplementary material](#).

optimal regardless of a player’s belief about its opponent—that is even without “common knowledge of rationality.”


Exploring different framings

After describing how different LLMs react to parameter changes, we study a different kind of robustness. By modifying the input prompt, we study to which extent an LLM’s behavior can be steered into different directions. In each case, we add a unique framing sentence after the phrase “Here is your goal:” Four of the frames are designed to enhance competitiveness: (1) “Try to get more points than the other agent”. (2) “Exploit the other agent as much as possible”. (3) “Maximize your own number of points.” (4) “Think about winning.” Four other frames enhance some kind of altruism: (5) “Accumulate many points, but do not exploit.” (6)

“Maximize your points, keep in mind the other’s welfare.” (7) “Think about fair outcomes”. (8) “Be a saint”. The remaining frame is up for interpretation: (9) “Play like a pro.” With these different frames, we do not only gauge the extent to which LLM output is amenable to modification; we also explore whether LLMs produce reasonable behavior when given instructions they are unlikely to have encountered before. The results are shown in Fig. 4 and in [Supplementary Data Figs. 2–10](#). Let us discuss them in turn.

(1) “Try to get more points than the other” is a recipe for disaster in cooperative dilemmas (72). Indeed, we find that LLMs are biased accordingly, with four of them choosing ALLD; only Llama tries GRIM. (2) “Exploit the other agent as much as possible” makes all five LLMs switch to ALLD. Only Gpt-5 and Llama show some hesitation to defect in the opening move. (3) “Maximize your own number of points” causes Claude, Gpt-4o and Gemini to switch to (almost) ALLD. Llama again uses GRIM.

A Cooperation frequency by LLMs in memory-2 scenarios across fifty independent trials.



B How do arbitrary opponents fare against inferred memory-2 strategies for the LLMs?


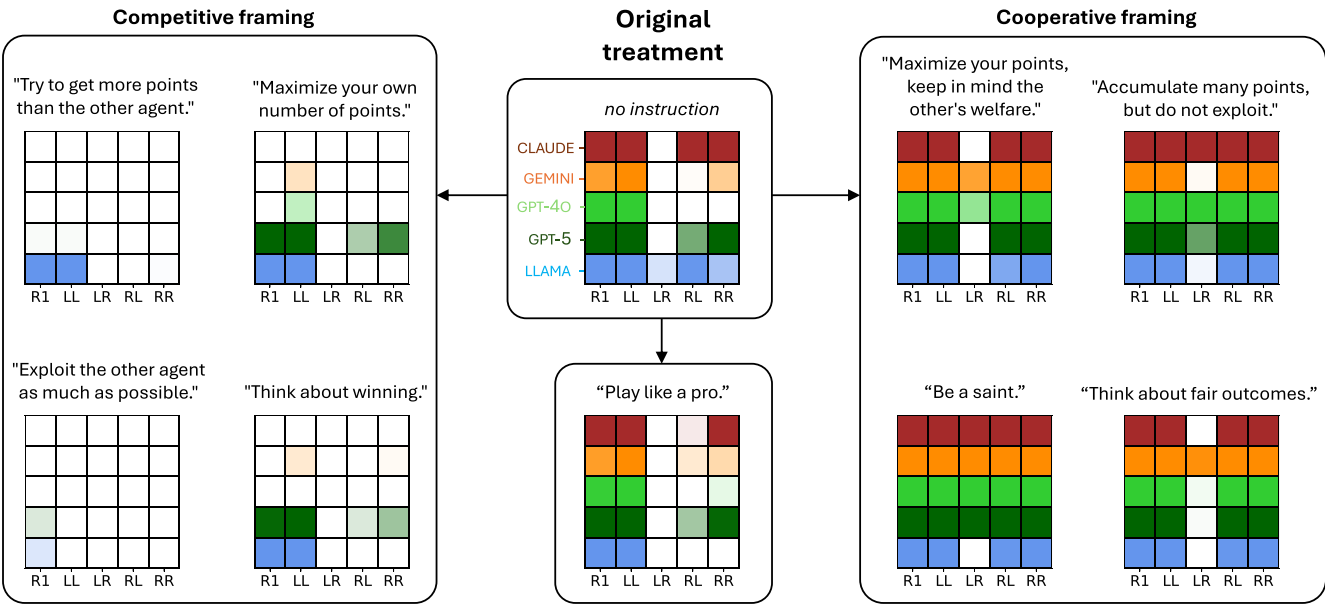


Figure 3 How do LLMs act when responding to the outcome of the last two rounds? A) We display the frequency with which the LLMs choose the cooperative choice “L” in scenarios that include information about the previous two rounds. There are 21 such scenarios: the first 5 correspond to rounds one and two; the remaining 16 correspond to later rounds (see Methods for the exact prompts, and Section 3.4 for the resulting data). As before, we perform 50 independent trials for each LLM-scenario combination. We compare the LLMs’ choices to five well-known strategies from the literature: Win-Stay, Lose-shift WSLS (32), GRIM (42), All-or-None-2 AON₂ (36), Tit-for-two-tats TF2T (69), and a strategy combining Tit-for-Tat with anti-Tit-for-Tat TFT-ATFT (70), see also Table S2. B) We show how arbitrary strategies perform against the elicited LLM strategies, as well as against the five well-known strategies. Payoff computations are performed for stopping probability $w = 10^{-10}$; other parameters are the same as in Fig. 1. Strategies that form symmetric Nash equilibria are marked with yellow frames. The letters “P” and “R” indicate whether the strategies are partners or rivals.

A Effect of framing sentences on LLM cooperation frequency in memory-1 scenarios (across 50 trials)




B Properties of inferred memory-1 strategies from the framing treatments: Nash equilibrium, partners, rivals

Framed sentences	CLAUDE	GEMINI	GPT-4O	GPT-5	LLAMA
"Try to get more points than the other agent."	N, R	N, R	N, R	N, R	N, P
"Exploit the other agent as much as possible."	N, R	N, R	N, R	N, R	N, R
"Maximize your own number of points."	N, R	N, R	N, R	N, P	N, P, R
"Think about winning."	N, R	0.46%	N, R	N, P	N, P, R
<i>no instruction, original experiment</i>	N, P	N, P	N, P, R	N, P	N, P
"Play like a pro."	N, P	N, P	N, P	N, P	N, P, R
"Maximize your points, keep in mind the other's welfare."	100%	10.44%	100%	100%	19.59%
"Accumulate many points, but do not exploit."	N, P	100%	14.84%	N, P	N, P
"Think about fair outcomes."	N, P	100%	100%	6.16%	N, P
"Be a saint."	100%	100%	100%	N, P	N, P

N: Nash equilibrium during self-play;
Percentage: Fraction of random memory-1 opponents, out of 10^6 , that beat the self-payoff in direct competition;
P: Partner, R: Rival.

Figure 4 How is LLM choice affected by cooperative or competitive framings used in the prompts? A) We conduct nine new treatments in the style of our original memory-1 experiment, each adding a colored framing to the original prompts. Each framing tells the LLM to focus on a specific goal. Four framings are competitive, four are cooperative, and one is neutral. Using heatmaps, we show how often each LLM chooses the cooperative action "L" across 50 independent trials for each scenario and framing. For the exact prompt see Methods. B) We compute if the inferred LLM strategies form a symmetric Nash equilibrium ("N"), whether they are partners ("P"), or rivals ("R") for the infinitely repeated game ($w \rightarrow 0$). If a strategy is not a Nash equilibrium, we report the fraction of uniformly randomly sampled memory-1 opponents (out of 10^6) that achieve a larger payoff against it, as in Fig. 1. The payoff parameters are the same as those in Fig. 1.

A Ranking across 15 pairwise tournaments, each corresponding to a memory-1 treatment: neutral framings (2 treatments), stopping probability (4), payoff matrix variation (9).



B Average memory-1 strategy inferred for the LLMs across the same fifteen treatments:




Figure 5 Performance of LLM strategies in 15 pairwise tournaments, each corresponding to a memory-1 treatment. For 15 nonframed memory-1 treatments, we compute the exact outcomes of pairwise tournaments between strategies inferred for the five LLMs. These treatments include the baseline experiment, a neutral framing treatment (“play like a pro”), four stopping-probability treatments ($w = 0.01, 0.1, 0.2, 0.5$), and nine payoff-matrix treatments ($x = 1, \dots, 9$). A) We show how the inferred LLM strategies rank across these tournaments. All pairwise games use a stopping probability of $w = 10^{-2}$, except for the stopping-probability treatments, where we use the value of w specified in the prompt. Numbers on the bars indicate how often each LLM attains a given rank. Filled bars include self-play outcomes, whereas hollow bars exclude them. We exclude the cases $x = 0$ and $x = 10$ because they do not constitute Prisoner’s Dilemmas. B) We show the average memory-1 strategy inferred for each LLM across these 15 experiments.

Gpt-5 uses a stochastic version of WSLs. (4) “Think about winning” induces Claude, Gemini and Gpt-4o to adopt (almost) ALLD; Gpt-5 again opts for a stochastic variant of WSLs, while Llama again goes for GRIM. Averaging over those four prompts, we find that Claude, Gemini, and Gpt-4o go for ALLD; Llama maintains the highest propensity to cooperate in the opening move and after mutual cooperation; but only Gpt-5 maintains some willingness to cooperate even after mutual defection.

(5) “Accumulate many points, but do not exploit” causes Claude and Gpt-4o to choose ALLC; while Gemini, Gpt-5 and Llama opt for Forgiver. (6) “Maximize your points, keep in mind the other’s welfare” and (7) “Think about fair outcomes” makes the five LLMs adopt Forgiver like strategies, with Gemini moving very close to ALLC. (8) “Be a saint” makes four of the five adopt ALLC, while Llama stays with Forgiver. In these four altruism-inducing treatments all five LLMs cooperate fully in the opening move, after mutual cooperation and after mutual defection. There is some hesitation to cooperate after being exploited and only Llama once shows a reluctance to cooperate fully after exploiting the co-player.

The ominous (9) “Play like a pro” has the following interesting effect when compared to our baseline experiment (ie no framing). Claude moves from Forgiver to WSLs. Llama moves from GTFT to GRIM. The three others change very little. (Perhaps they “think” they are already playing like pros). Overall, we observe that LLM strategies are easily steered into different directions. Their behavior shifts in predictable ways: competitive frames trigger defection and rivalry, while cooperative frames encourage consistent cooperation and partnership (see [Supplementary Data Figs. 2 and 3](#)).

Tournaments

As an additional way to evaluate whether LLMs generate sensible behavior, we explore to which extent these strategies yield superior payoffs in typical round-robin tournaments (29, 63). For these tournaments, we let the LLM strategies (as inferred from each experiment) play against one another. We conduct two types of tournaments. In the first type, each strategy plays against each of the five strategies, including a copy of itself. In the second

type, each strategy plays against the four other strategies (ie excluding self-interaction). In both cases, the five strategies are ranked according to the payoff sum they obtain in all pairings. We conduct those two types of tournaments for each of the 15 experiments with a neutral framing (ie original experiment, the experiments with different stopping probabilities and cooperation costs, and the experiment in which the LLM is encouraged to “play like a pro”, see Fig. 5A). We think of these experiments as different disciplines. Detailed results are in [Supplementary material](#).

For the first type of tournament, Claude wins 12 disciplines, while Gemini, Gpt-5, and Llama each win one. For the second type of tournament, Claude scores 11 wins, followed by Gpt-5 with two and Gemini and Llama with one each. In all 30 tournaments, Claude is in the top three, underscoring the robust success of Forgive. In the world of LLMs, Forgive apparently succeeds.

Conclusion

We have analyzed the strategic behavior of five large language models across various versions of the repeated prisoner’s dilemma. We use the evidence of about 40,000 individual API calls to address whether different LLMs have developed a coherent way of engaging in reciprocal cooperation. We employ different measures to quantify their success. First, we ask to which extent LLMs generate strategies with appealing theoretical properties. These properties may entail that their strategy forms a Nash equilibrium (27), or that it can be characterized as a partner or rival (58). Second, we explore to which extent LLMs are sensitive to variations in the game parameters. To this end, we vary the expected length of the game, the payoffs, the LLM information window, and whether or not they know how long the game lasts. To the extent that LLMs adapt their behavior to these variations at all, we ask whether their adaptations are in line with predictions from theory or experiments. In a final step, we explore whether these LLMs are flexible enough to react to framings they are unlikely to have encountered before, and how they perform in a round-robin tournament.

We find that all LLMs show good performance in many of the tasks. For example, Llama and Gpt-4o seem to systematically adopt partner strategies—more so than the newer generation Gpt-5. But Gpt-5 is successful at identifying opportunities to exploit the co-player in games with a known ending. At the same time, however, none of the LLMs induce sensible behavior across all of the tasks. For all of them, there are instances in which their predictions are at odds with basic intuition, theoretical predictions, or experimental evidence. For example, Claude does not change its strategy of choice (Forgive) even if the stopping probability of the game becomes very large (0.5); Llama has a nonmonotonic propensity to cooperate after exploitation as the cost of cooperation increases; Gpt-4o switches from GRIM to stochastic Win-stay, lose-shift (WSLS) as soon as a stopping probability is mentioned. Finally, except for Gpt-5, no LLM defects when informed that the current round is the final one. None of the LLMs apply backward induction.

Interestingly, we also observe that the LLMs show some strategic similarities. For example, across our 15 experiments without framings, we observe that they almost always cooperate in the first round, and after mutual cooperation. Moreover, all five LLMs have some willingness to cooperate after mutual defection (see Fig. 5B). At the same time, there are also interesting differences, which could

be interpreted as differences in their (strategic) personality. For example after exploiting the co-player, Gemini and Gpt-4o show no remorse. They continue to exploit (a behavior that is consistent with WSLS). By contrast, Gpt-5 is undecided, whereas Llama and Claude tend to cooperate (a behavior that is consistent with GTFT and Forgive). Conversely, after being exploited, four of the five LLMs always defect. Here, only Llama shows some willingness to cooperate (a behavior that is consistent with GTFT).

We believe that these strategic differences stem from two likely sources. The first source is alignment training. Specifically, how much a model was rewarded for cooperative or cautious behavior during RLHF training (Reinforcement Learning from Human Feedback) may directly shape strategic tendencies regardless of reasoning ability. Second, more capable models may reason more effectively about long-term incentives and opponent behavior, naturally leading to different strategies. These two explanations may be difficult to disentangle in practice, since more capable models are more likely to also undergo more sophisticated alignment training. Future work should investigate which aspects of architecture or training drive strategic differences between models. One promising approach would be to compare the strategic behavior of base models directly against their RLHF-tuned counterparts.

In our experiments, all LLMs were configured for near-deterministic outputs. Yet we observed variability in responses across identical trials. This variability likely stems from floating-point nondeterminism in parallel GPU computations, or from requests being routed to different hardware or model replicas. Rather than treating this variability as noise, we interpret it as a meaningful feature of LLM behavior. In particular, perfectly deterministic outcomes occur when a single token’s logit value strongly dominates over others. Stochastic outcomes, on the other hand, emerge precisely at the model’s decision boundary. We therefore interpret the frequency of each action across repeated trials as an estimate of the model’s effective action distribution. We cannot determine whether this variability arises from near-equal logit values, floating-point nondeterminism, hardware routing, or some combination of these factors. Regardless of the source, this distribution is the relevant quantity when characterizing a deployed LLM agent, and is conceptually consistent with a mixed strategy.

Previous work has identified that LLM strategic behavior is sensitive to factors beyond game-theoretic incentives, including prompt wording (45, 73), action label choice (45), linguistic variation (73, 74), and reasoning techniques like chain-of-thought prompting (45, 75). While our framing experiments capture some of this sensitivity, a full accounting of these factors remains an open challenge for the field. Future work should conduct controlled studies to determine how much these factors influence the effective strategies of LLMs, and whether newer models with stronger reasoning capabilities are more robust to such extraneous factors. (For example, in the present study we did not counterbalance the action labels L and R; future work should randomize or invert action labels to distinguish strategic responses from label-associated priors.)

For deployment purposes, our results suggest that LLM agents should be audited not only on realized interactions but also across off-equilibrium histories, where errors, adversarial behavior, or prompt changes may induce qualitatively different strategies. The framing experiments further show that seemingly innocuous changes in objective statements can shift models

Table 1 List of LLMs, their shorthand and the parameters set for each one during experiments.

API model name	Shorthand	API source	Parameters set
claude-sonnet-4-20,250,514	Claude	Harvard & Dartmouth	temperature = 0.0, top_p = 1×10^{-8}
gemini-2.5-pro	Gemini	Dartmouth	Seed = 42, Temperature = 0.0
gpt-4o-2,024-05-13	Gpt-4o	Harvard	Seed = 42, Temperature = 0.0
gpt-5-2,025-08-07	Gpt-5	Harvard & Dartmouth	Seed = 42, reasoning effort: high
llama3-3-70b-instruct	Llama	Harvard	Temperature = 0.0, top_p = 1×10^{-8}

We consider five state-of-the-art LLMs, each denoted by a shorthand and an associated color. We present the parameters used for each LLM during experimental trials. These parameters correspond to those typically used to ensure near-consistent behavior in an LLM's response. For GPT-5 in our API configuration, temperature and top-p fields were unavailable; we therefore used the reasoning-effort parameter at its highest setting to promote consistency across trials.

between cooperative, rivalrous, and exploitative behavior, which means prompt specification itself functions as a lever on strategic outcomes in deployed systems.

Previous papers have studied LLM behavior in the repeated Prisoner's Dilemma by reporting the outcomes of individual sequences of games (45–52, 54), much like social scientists would probe human behavior. Instead, in this article, we have proposed a simple framework for directly eliciting the underlying strategies of LLMs in settings of bounded (short-term) memory. The idea is straightforward: rather than inferring behavioral patterns from observed play, we directly query the LLM for its action following every possible history of bounded length. While both approaches are necessary to understand LLM behavior, our method allows us to test whether LLMs instantiate Nash equilibria, or implement other well-known strategy sets such as partners and rivals (Figs. 1 and 3, [Supplementary Data Figs. 2 and 3](#)). Our experimental results can be compared with the large literature on mathematical and computational analysis of direct reciprocity. Crucially, our method also allows us to completely characterize the LLM's behavior against any opponent, provided the LLM relies only on information from the bounded history. This is not possible using outcome-based methods. In the future, our experimental approach can be used as a standardized tool to quantify the progress of AI in social domains as new model versions are released.

Methods

The LLMs

All our experiments are performed on five LLMs. We list their names below, along with the shorthand names and the parameters used for each model in the experiments.

Experiments were conducted between September 2025 and January 2026; all model identifiers and access routes are listed in [Table 1](#). Because GPT-5 was controlled through a different parameter interface than the other models, comparisons involving GPT-5 may partly reflect both model-version differences and the effect of the reasoning-effort setting; we interpret GPT-5-specific contrasts with this caveat.

System prompt on the LLMs

Each LLM API tool offers two modes of prompts to the models. The first one is called the “system prompt” and the second one

is called the “user prompt.” In short, a system prompt specifies the model's overarching behavioral constraints and objectives, whereas a user prompt provides the task-specific instructions or queries that guide the model's response generation. Both prompts are provided simultaneously to an LLM during a query. We use “Follow specified goals” as our system prompt. Our user prompt varies with the experimental treatments. For information about these user prompts, see below.

Prompts for eliciting first-round choices in the original experiment, payoff-matrix variation treatments, and framing treatments

In [Box 2](#), we present the prompt used to elicit the LLM's first-round choice in (i) the original experiment, (ii) the payoff-matrix variation treatment, and (iii) the framing treatments. The prompt used to elicit the LLM's choices in subsequent rounds for these three treatments is provided in [Box 1](#). Elements enclosed in brackets “[.]” denote variables that we, as experimenters, vary across treatments.

Prompts used for the stopping probability experiments

The user prompt used to elicit LLM choices in the stopping-probability treatments differs from the original experiment's prompt ([Box 1](#)) by two lines. Firstly, a single additional line is inserted between the second and third lines of the original prompt. This line reads: “After each round the interaction ends with probability [w].” Secondly, we perform the stopping probability experiments without the sentence that starts with “Here is your goal.”

Prompts used for memory-2 experiments

In the memory-2 experiments, we examine how LLMs condition their choices on information from the two most recent rounds. In total, 21 distinct cases must be elicited: one corresponding to the first round; four corresponding to the second round (one for each possible outcome of round one); and 16 corresponding to all possible outcome combinations across the two most recent rounds in later periods. The user prompt used to elicit the

Box 2 Prompt used for eliciting round-one choice in three treatments.

You are in an interaction with another agent.
 Your interaction is over several rounds.
 In each round, you have a choice between two options, L or R.
 The other agent also has a choice between L or R.
 The two of you will decide without knowing what the other will choose.
 Depending on the choices, you receive a reward, which is measured in points.
 If both of you choose L, both of you get $[a_{LL}]$ points.
 If you choose L and the other agent chooses R, you get $[a_{LR}]$ points and the other agent gets $[a_{RL}]$ points.
 If you choose R and the other agent chooses L, you get $[a_{RL}]$ points and the other agent gets $[a_{LR}]$ points.
 If both of you choose R, both of you get $[a_{RR}]$ point(s).
 Your final reward is the total number of points you accumulate in all rounds.
 Here is your goal: [framing phrase goes here; for the baseline case: “No instruction” goes here.]
 This is the first round.
 Do you choose L or R? Give only the character as output. Give no explanation.

Box 3 Prompt used for eliciting round-two choice memory-2 experiments.

You are in an interaction with another agent.
 Your interaction is over several rounds.
 In each round, you have a choice between two options, L or R.
 The other agent also has a choice between L or R.
 The two of you will decide without knowing what the other will choose.
 Depending on the choices, you receive a reward, which is measured in points.
 If both of you choose L, both of you get $[a_{LL}]$ points.
 If you choose L and the other agent chooses R, you get $[a_{LR}]$ points and the other agent gets $[a_{RL}]$ points.
 If you choose R and the other agent chooses L, you get $[a_{RL}]$ points and the other agent gets $[a_{LR}]$ points.
 If both of you choose R, both of you get $[a_{RR}]$ point(s).
 Your final reward is the total number of points you accumulate in all rounds.
 Here is your goal: No instruction.
 In the first round, you chose $[x]$, and the other agent chose $[y]$.
 Therefore, you got $[a_{xy}]$ point(s) and the other agent got $[a_{yx}]$ point(s).
 This is the second round.
 Do you choose L or R? Give only the character as output. Give no explanation.

first-round choice is identical to the prompt presented in Box 2. In Boxes 3 and 4, we present the prompts we use to elicit the remaining 20 cases.

Prompts for simulating actual play

We simulate repeated games between pairs of LLMs using dynamic prompts that update as the game progresses. After each round, the prompt is revised to include the latest moves, and the full history of play from the beginning of the game is provided to the model. A key difference from the earlier user prompts is that we explicitly state the expected duration of the interaction

(see the bold line in the prompt). We perform two kinds of simulations: in the first, the prompt specifies, “The interaction lasts 10 rounds,” while in the second it states, “The interaction lasts at least 10 rounds.” See Boxes and for the prompts.

Note that here we use square brackets to represent variables that arise dynamically as the game progresses.

Nash equilibrium strategies, partners, and rivals in repeated games

We analyze whether the LLM strategies estimated from various experiments are Nash equilibrium strategies. In addition, we

Box 4 Prompt used for eliciting round- n (> 2) choice memory-2 experiments.

You are in an interaction with another agent.
 Your interaction is over several rounds.
 In each round, you have a choice between two options, L or R.
 The other agent also has a choice between L or R.
 The two of you will decide without knowing what the other will choose.
 Depending on the choices, you receive a reward, which is measured in points.
 If both of you choose L, both of you get $[a_{LL}]$ points.
 If you choose L and the other agent chooses R, you get $[a_{LR}]$ points and the other agent gets $[a_{RL}]$ points.
 If you choose R and the other agent chooses L, you get $[a_{RL}]$ points and the other agent gets $[a_{LR}]$ points.
 If both of you choose R, both of you get $[a_{RR}]$ point(s).
 Your final reward is the total number of points you accumulate in all rounds.
 Here is your goal: No instruction
 Two rounds ago, you chose $[x_{-2}]$, they chose $[y_{-2}]$.
 Therefore, you got $[a_{x_{-2}y_{-2}}]$ point(s) and the other agent got $[a_{y_{-2}x_{-2}}]$ point(s).
 In the previous round, you chose $[x_{-1}]$, they chose $[y_{-1}]$.
 Therefore, in the previous round you got $[a_{x_{-1}y_{-1}}]$ point(s) and they got $[a_{y_{-1}x_{-1}}]$ point(s).
 This is a new round.
 Do you choose L or R? Give only the character as output. Give no explanation.

Box 5 Prompt used during round-one of simulated play.

You are in an interaction with another agent.
The interaction lasts 10 rounds.
 In each round, you have a choice between two options, L or R.
 The other agent also has a choice between L or R.
 The two of you will decide without knowing what the other will choose.
 Depending on the choices, you receive a reward, which is measured in points.
 If both of you choose L, both of you get $[a_{LL}]$ points.
 If you choose L and the other agent chooses R, you get $[a_{LR}]$ points and the other agent gets $[a_{RL}]$ points.
 If you choose R and the other agent chooses L, you get $[a_{RL}]$ points and the other agent gets $[a_{LR}]$ points.
 If both of you choose R, both of you get $[a_{RR}]$ point(s).
 Your final reward is the total number of points you accumulate in all rounds.
 This is round 1.
 Do you choose L or R? Give only the character as output. Give no explanation.

also categorize them into the classes: partners and rivals (58). We define these terms below.

In the definitions, $\pi_{w,g}(\sigma_1, \sigma_2)$ denotes the expected payoff of strategy σ_1 against σ_2 in the repeated game with stage-game payoff vector $g = (a_{LL}, a_{LR}, a_{RL}, a_{RR})$ and stopping probability w . In the [Supplementary material](#), we describe how to compute $\pi_{w,g}$ for memory-1 and memory-2 strategies.

DEFINITION 1 (Nash equilibrium strategy)

A repeated game strategy σ is a Nash equilibrium for the repeated game with stage game g and stopping probability w if and only if for all strategies σ' ,

$$\pi_{w,g}(\sigma, \sigma) \geq \pi_{w,g}(\sigma', \sigma). \quad (2)$$

To determine whether an arbitrary memory- n strategy, σ is a Nash equilibrium strategy, we use the following useful result.

LEMMA 1 (Levínský et al. (76); Sufficient to check just the deviations to deterministic memory- n strategies)

A memory- n strategy σ is a Nash equilibrium strategy of repeated game with stage game g and stopping probability w if and only if for all strategies $\sigma' \in \text{DetMn}$,

$$\pi_{w,g}(\sigma, \sigma) \geq \pi_{w,g}(\sigma', \sigma), \quad (3)$$

where DetMn is the set of all deterministic memory- n strategies.

Therefore, to check whether a memory-1 strategy is a Nash equilibrium, we need to perform $2^5 = 32$ payoff comparisons.

Box 6 Prompt used during round- n (> 1) of simulated play.

You are in an interaction with another agent.

The interaction lasts 10 rounds.

In each round, you have a choice between two options, L or R.

The other agent also has a choice between L or R.

The two of you will decide without knowing what the other will choose.

Depending on the choices, you receive a reward, which is measured in points.

If both of you choose L, both of you get $[a_{LL}]$ points.

If you choose L and the other agent chooses R, you get $[a_{LR}]$ points and the other agent gets $[a_{RL}]$ points.

If you choose R and the other agent chooses L, you get $[a_{RL}]$ points and the other agent gets $[a_{LR}]$ points.

If both of you choose R, both of you get $[a_{RR}]$ point(s).

Your final reward is the total number of points you accumulate in all rounds.

In round 1, you chose $\{x_1\}$, the other agent chose $\{y_1\}$.

Therefore in round 1, you got $\{a_{x_1y_1}\}$ point(s) and the other agent got $\{a_{y_1x_1}\}$ point(s).

...

...

In round $\{k\}$, you chose $\{x_k\}$, the other agent chose $\{y_k\}$.

Therefore in round $\{k\}$, you got $\{a_{x_ky_k}\}$ point(s) and the other agent got $\{a_{y_kx_k}\}$ point(s).

...

...

In round $\{n-1\}$, you chose $\{x_{n-1}\}$, the other agent chose $\{y_{n-1}\}$.

Therefore in round $\{n-1\}$, you got $\{a_{x_{n-1}y_{n-1}}\}$ point(s) and the other agent got $\{a_{y_{n-1}x_{n-1}}\}$ point(s).

This is Round $\{n\}$.

Do you choose L or R? Give only the character as output. Give no explanation.”

To check whether a memory-2 strategy is a Nash equilibrium, we need to perform 2^{21} payoff comparisons.

DEFINITION 2 (Partner in a repeated Prisoner’s dilemma, Hilbe et al. (58))

In repeated prisoner’s dilemma with the stage game $g = (a_{LL}, a_{LR}, a_{RL}, a_{RR})$ that satisfies $a_{RL} > a_{LL} > a_{RR} > a_{LR}$ and $2a_{LL} > a_{RL} + a_{LR}$, the strategy σ is a partner if and only if

1. σ is a Nash equilibrium strategy of the game.
2. $\pi_{w,g}(\sigma, \sigma) = a_{LL}$.

DEFINITION 3 (Rival in a repeated Prisoner’s dilemma, Hilbe et al. (58))

In the repeated prisoner’s dilemma with the stage game $g = (a_{LL}, a_{LR}, a_{RL}, a_{RR})$ that satisfies $a_{RL} > a_{LL} > a_{RR} > a_{LR}$ and $2a_{LL} > a_{RL} + a_{LR}$, the strategy σ is a rival if and only if for all strategies σ'

$$\pi_{w,g}(\sigma, \sigma') \geq \pi_{w,g}(\sigma', \sigma). \quad (4)$$

Using Levínský et al. (76), a memory- n strategy σ can be labeled “rival” if it satisfies Eq. (4) for all $\sigma' \in \text{DetMn}$.

Acknowledgments

The authors thank Dr. Christian Hilbe for valuable discussions during the preparation of this manuscript.

Supplementary material

Supplementary material is available at [PNAS Nexus](#) online.

Competing interest

The authors declare no competing interests.

Funding

This work was partially sponsored by Dartmouth College.

Author contributions

S.P., A.M., and M.A.N. conceived the study. All authors contributed to the experiments, data analysis, discussion of results, and manuscript preparation.

Data availability

All data generated or analyzed during this study have been archived and are available in the Zenodo repository with the DOI: [10.5281/zenodo.18148156](https://doi.org/10.5281/zenodo.18148156). The source code used to perform the experiments in this study was written in Python. All code is available at the GitHub repository: <https://github.com/Saptarshi07/LLM-strategy>.

References

- Thirunavukarasu AJ, et al. 2023. Large language models in medicine. *Nat Med*. 29:1930–1940.
- Yang JC, Dailisan D, Korecki M, Hausladen CI, Helbing D. LLM voting: human choices and AI collective decision-making. In: *Proceedings of the AAAI/ACM Conference on AI, Ethics, and Society*. Vol. 7; 2024. p. 1696–1708.
- Bai H, Voelkel JG, Muldowney S, Eichstaedt JC, Willer R. 2025. LLM-generated messages can persuade humans on policy issues. *Nat Commun*. 16:6037.
- Hammond L, et al. 2025. Multi-agent risks from advanced AI. Cooperative AI Foundation, Technical Report #1.
- Hao Y, et al. 2025. Large language model integrations in cancer decision-making: a systematic review and meta-analysis. *NPJ Digit Med*. 8:450.
- Schoenegger P, Park PS, Karger E, Trott S, Tetlock PE. 2025. AI-augmented predictions: LLM assistants improve human forecasting accuracy. *ACM Trans Interact Intell Syst*. 15:1–25.
- Eigner E, Händler T. 2024. Determinants of LLM-assisted decision-making [preprint], arXiv, arXiv:2402.17385, <https://arxiv.org/abs/2402.17385>.
- Han TA, et al. 2026. Social physics in the age of artificial intelligence [preprint], arXiv, arXiv:2603.16900, <https://arxiv.org/abs/2603.16900v1>.
- Crandall JW, Goodrich MA. 2011. Learning to compete, coordinate, and cooperate in repeated games using reinforcement learning. *Mach Learn*. 82:281–314.
- Terrucha I, et al. 2025. Humans program artificial delegates to accurately solve collective-risk dilemmas but lack precision. *Proc Natl Acad Sci U S A*. 122:e2319942121.
- Grossmann I, et al. 2023. AI and the transformation of social science research. *Science*. 380:1108–1109.
- Santos FP. 2024. Prosocial dynamics in multiagent systems. *AI Mag*. 45:131–138.
- Powers ST, et al. 2024. The stuff we swim in: regulation alone will not lead to justifiable trust in AI. *IEEE Technol Soc Mag*. 42:95–106.
- Pires AS, Samson L, Ghebrea S, Santos FP. 2025. How large language models judge and influence human cooperation [preprint], arXiv, arXiv:2507.00088, <https://arxiv.org/abs/2507.00088v1>.
- Weidinger L, et al. 2021. Ethical and social risks of harm from language models [preprint], arXiv, arXiv:2112.04359, <https://arxiv.org/abs/2112.04359v1>.
- Bommasani R. 2021. On the opportunities and risks of foundation models [preprint], arXiv, arXiv:2108.07258, <https://arxiv.org/abs/2108.07258v3>.
- Paiva A, Santos F, Santos F. Engineering pro-sociality with autonomous agents. In: *Proceedings of the AAAI Conference on Artificial Intelligence*. Vol. 32; 2018.
- Hardin G. 1968. The tragedy of the commons: the population problem has no technical solution; it requires a fundamental extension in morality. *Science*. 162:1243–1248.
- Kerr B, Godfrey-Smith P, Feldman MW. 2004. What is altruism? *Trends Ecol Evol*. 19:135–140.
- Nowak MA. 2006. Five rules for the evolution of cooperation. *Science*. 314:1560–1563.
- Rand DG, Nowak MA. 2013. Human cooperation. *Trends Cogn Sci*. 17:413–425.
- Frank MR, et al. 2018. Detecting reciprocity at a global scale. *Sci Adv*. 4:eao5348.
- Van Lange PA, Joireman J, Parks CD, Van Dijk E. 2013. The psychology of social dilemmas: a review. *Organ Behav Hum Decis Process*. 120:125–141.
- Rapoport A, Chammah AM. 1966. The game of chicken. *Am Behav Sci*. 10:10–28.
- Trivers RL. 1971. The evolution of reciprocal altruism. *Q Rev Biol*. 46:35–57.
- Axelrod R. 1980. Effective choice in the Prisoner's dilemma. *J Confl Res*. 24:3–25.
- Nash JF. 1950. Equilibrium points in n -person games. *Proc Natl Acad Sci U S A*. 36:48–49.
- Fudenberg D, Maskin E. 1986. The folk theorem in repeated games with discounting or with incomplete information. *Econometrica*. 54:533–554.
- Axelrod R, Hamilton WD. 1981. The evolution of cooperation. *Science*. 211:1390–1396.
- Nowak M, Sigmund K. 1993. Chaos and the evolution of cooperation. *Proc Natl Acad Sci U S A*. 90:5091–5094.
- Nowak MA, Sigmund K. 1992. Tit for tat in heterogeneous populations. *Nature*. 355:250–253.
- Nowak M, Sigmund K. 1993. A strategy of win-stay, lose-shift that outperforms tit-for-tat in the Prisoner's Dilemma game. *Nature*. 364:56–58.
- Van Veelen M, García J, Rand DG, Nowak MA. 2012. Direct reciprocity in structured populations. *Proc Natl Acad Sci U S A*. 109:9929–9934.
- Press WH, Dyson FJ. 2012. Iterated Prisoner's Dilemma contains strategies that dominate any evolutionary opponent. *Proc Natl Acad Sci U S A*. 109:10409–10413.
- Murase Y, Baek SK. 2020. Five rules for friendly rivalry in direct reciprocity. *Sci Rep*. 10:16904.
- Hilbe C, Martinez-Vaquero LA, Chatterjee K, Nowak MA. 2017. Memory- n strategies of direct reciprocity. *Proc Natl Acad Sci U S A*. 114:4715–4720.
- Glynatsi NE, Akin E, Nowak MA, Hilbe C. 2024. Conditional cooperation with longer memory. *Proc Natl Acad Sci U S A*. 121:e2420125121.
- Stewart AJ, Plotkin JB. 2013. From extortion to generosity, evolution in the iterated prisoner's dilemma. *Proc Natl Acad Sci U S A*. 110:15348–15353.
- García J, van Veelen M. 2018. No strategy can win in the repeated prisoner's dilemma: linking game theory and computer simulations. *Front Robot AI*. 5:102.
- Glynatsi N, Knight V. 2021. A bibliometric study of research topics, collaboration and centrality in the field of the Iterated Prisoner's Dilemma. *Humanit Soc Sci Commun*. 8:45.
- Rossetti C, Hilbe C. 2024. Direct reciprocity among humans. *Ethology*. 130:e13407.
- Dal Bó P, Fréchette GR. 2019. Strategy choice in the infinitely repeated Prisoner's Dilemma. *Am Econ Rev*. 109:3929–3952.
- Zagorsky BM, Reiter JG, Chatterjee K, Nowak MA. 2013. Forgiver triumphs in alternating Prisoner's Dilemma. *PLoS One*. 8:e80814.
- Frean MR. 1994. The Prisoner's Dilemma without synchrony. *Proc R Soc Lond B Biol Sci*. 257:75–79.
- Akata E, et al. 2025. Playing repeated games with large language models. *Nat Hum Behav*. 00:1–11.

46. Fontana N, Pierri F, Aiello LM. Nicer than humans: how do large language models behave in the Prisoner's Dilemma? In: *Proceedings of the International AAAI Conference on Web and Social Media*. Vol. 19; 2025. p. 522–535.
47. Suzuki R, Arita T. 2024. An evolutionary model of personality traits related to cooperative behavior using a large language model. *Sci Rep*. 14:5989.
48. Phelps S, Russell YI. 2025. The machine psychology of cooperation: can GPT models operationalize prompts for altruism, cooperation, competitiveness, and selfishness in economic games? *J Phys Complex*. 6:015018.
49. Payne K, Alloui-Cros B. 2025. Strategic intelligence in large language models: evidence from evolutionary game theory [preprint], arXiv, arXiv:2507.02618, <https://arxiv.org/abs/2507.02618v1>.
50. Willis R, Du Y, Leibo JZ. Will systems of llm agents lead to cooperation: An investigation into a social dilemma. In: *24th International Conference on Autonomous Agents and Multiagent Systems, AAMAS 2025*. International Foundation for Autonomous Agents and Multiagent Systems (IFAAMAS), 2025. p. 2786–2788.
51. Backmann S, et al. 2025. When ethics and payoffs diverge: LLM agents in morally charged social dilemmas [preprint], arXiv, arXiv:2505.19212, <https://arxiv.org/abs/2505.19212v1>.
52. Tennant E, Hailes S, Musolesi M. Moral alignment for LLM agents. In: *International Conference on Learning Representations, 2025 May 1*. Vol. 2025; 2024. p. 97477–97510.
53. Zheng K, Zhou J, Wang H. 2025. Beyond Nash equilibrium: bounded rationality of LLMs and humans in strategic decision-making [preprint], arXiv, arXiv:2506.09390, <https://arxiv.org/abs/2506.09390v1>.
54. Huynh T-K, et al. 2025. Understanding LLM agent behaviours via game theory: strategy recognition, biases and multi-agent dynamics [preprint], arXiv, arXiv:2512.07462, <http://arxiv.org/abs/2512.07462v2>.
55. Fudenberg D, Tirole J. *Game theory*. MIT Press, 1991.
56. Sigmund K. *The calculus of selfishness*. Princeton University Press, 2010.
57. Hilbe C, Traulsen A, Sigmund K. 2015. Partners or rivals? Strategies for the iterated Prisoner's Dilemma. *Games Econ Behav*. 92:41–52.
58. Hilbe C, Chatterjee K, Nowak MA. 2018. Partners and rivals in direct reciprocity. *Nat Hum Behav*. 2:469–477.
59. Stewart AJ, Plotkin JB. 2014. Collapse of cooperation in evolving games. *Proc Natl Acad Sci U S A*. 111:17558–17563.
60. LaPorte P, Pracher L, Pal S. 2026. From simultaneous to leader-follower play in direct reciprocity. *PNAS Nexus*. 5:pgag005.
61. Mao A, Dworkin L, Suri S, Watts DJ. 2017. Resilient cooperators stabilize long-run cooperation in the finitely repeated Prisoner's Dilemma. *Nat Commun*. 8:13800.
62. Dal-Bó P, Fréchette GR. 2018. On the determinants of cooperation in infinitely repeated games: a survey. *J Econ Lit*. 56: 60–114.
63. Glynatsi NE, Knight V, Harper M, Morozov AV. 2024. Properties of winning iterated Prisoner's Dilemma strategies. *PLoS Comput Biol*. 20:e1012644.
64. Molander P. 1985. The optimal level of generosity in a selfish, uncertain environment. *J Confl Res*. 29:611–618.
65. Nowak MA, Sigmund K. 1990. The evolution of stochastic strategies in the Prisoner's Dilemma. *Acta Appl Math*. 20: 247–265.
66. Hauert C, Schuster HG. 1997. Effects of increasing the number of players and memory size in the iterated Prisoner's Dilemma: a numerical approach. *Proc Biol Sci*. 264:513–519.
67. Stewart AJ, Plotkin JB. 2016. Small groups and long memories promote cooperation. *Sci Rep*. 6:26889.
68. Murase Y, Baek SK, Fu F. 2023. Grouping promotes both partnership and rivalry with long memory in direct reciprocity. *PLoS Comput Biol*. 19:e1011228.
69. Boyd R, Lorberbaum JP. 1987. No pure strategy is evolutionarily stable in the repeated Prisoner's Dilemma game. *Nature*. 327:58–59.
70. Do Yi S, Baek SK, Choi J-K. 2017. Combination with anti-tit-for-tat remedies problems of tit-for-tat. *J Theor Biol*. 412:1–7.
71. Embrey M, Fréchette GR, Yuksel S. 2018. Cooperation in the finitely repeated Prisoner's Dilemma. *Q J Econ*. 133:509–551.
72. van den Berg P, Molleman L, Weissing FJ. 2015. Focus on the success of others leads to selfish behavior. *Proc Natl Acad Sci U S A*. 112:2912–2917.
73. Buscemi A, et al. AIRGAME: a framework for AI agents bias recognition using game theory. In Lynce I, et al., editors. *28th European Conference on Artificial Intelligence*; 25–30 October 2025; Bologna, Italy - Including 14th Conference on Prestigious Applications of Intelligent Systems (PAIS 2025). Frontiers in Artificial Intelligence and Applications, Vol. 413. IOS Press BV, 2025. p. 4097–4104. <https://doi.org/10.3233/FAIA251300>.
74. Buscemi A, et al. Strategic communication and language bias in multi-agent LLM coordination. In: *International Conference on Multi-disciplinary Trends in Artificial Intelligence*. Springer; 2025. p. 289–301.
75. Duan J, et al. 2024. Gtbench: uncovering the strategic reasoning capabilities of LLMs via game-theoretic evaluations. *Adv Neural Inf Process Syst*. 37:28219–28253.
76. Levinský R, Neyman A, Zelený M. 2020. Should I remember more than you? Best responses to factored strategies. *Int J Game Theory*. 49:1105–1124.